GSM8K Static-K comparison | seed 42 | 30 episodes | learning rate 1e-5

Best-checkpoint held-out numeric EM (95% interval)

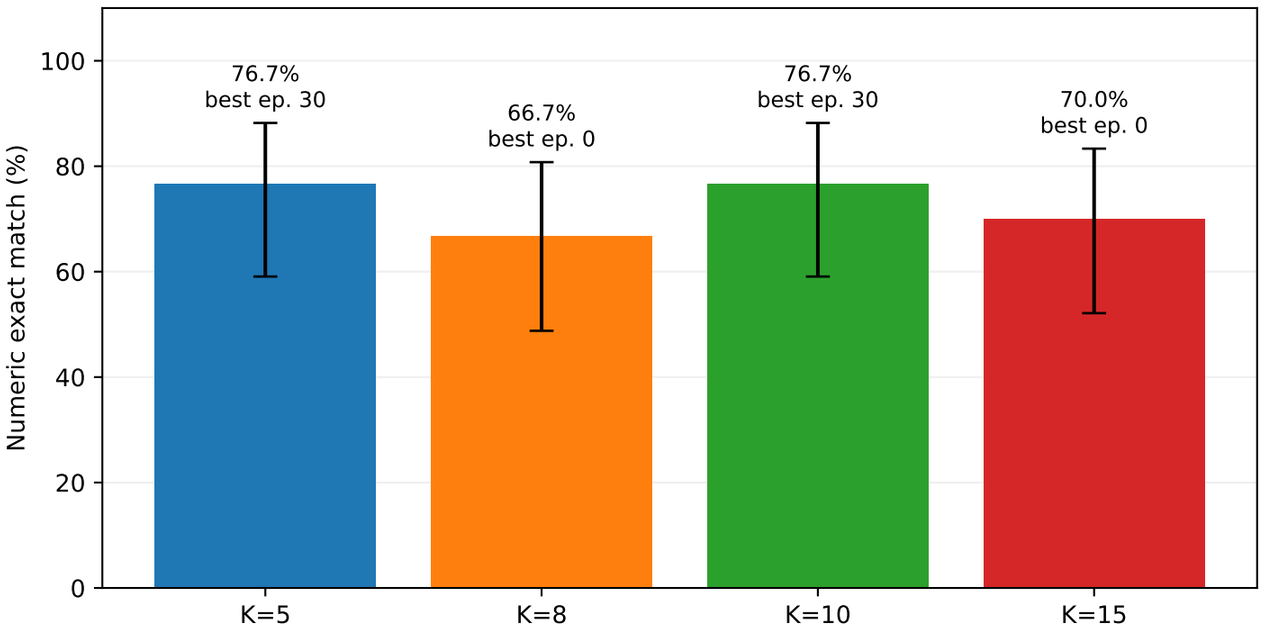

Validation EM at measured checkpoints

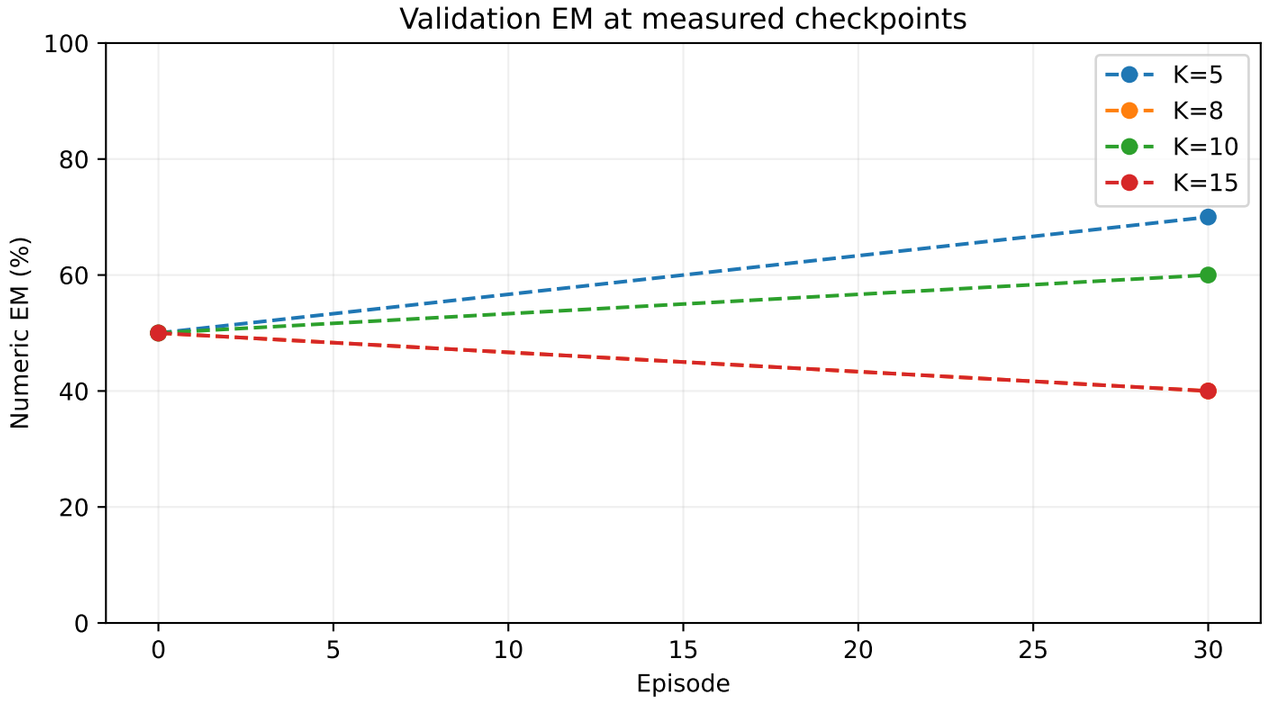

Training reward: trailing 5-episode mean

Distinct pipelines generated per episode

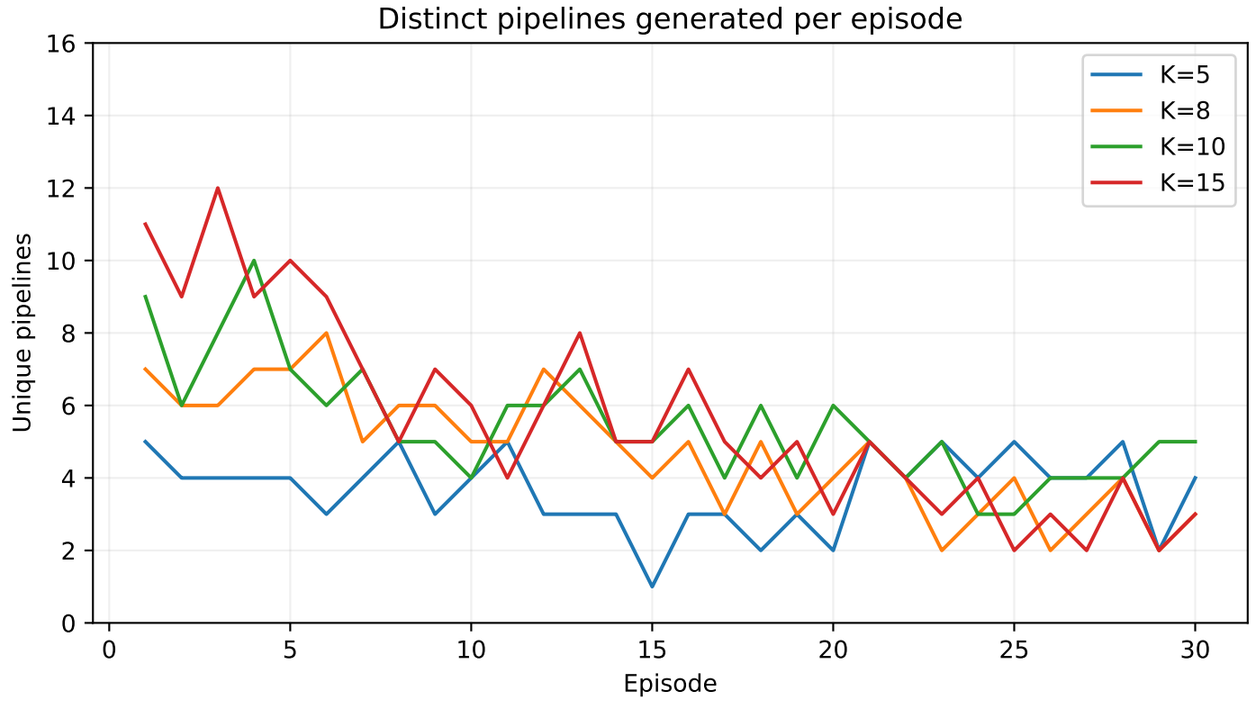

Gradient norm: original Figure 2 summed loss

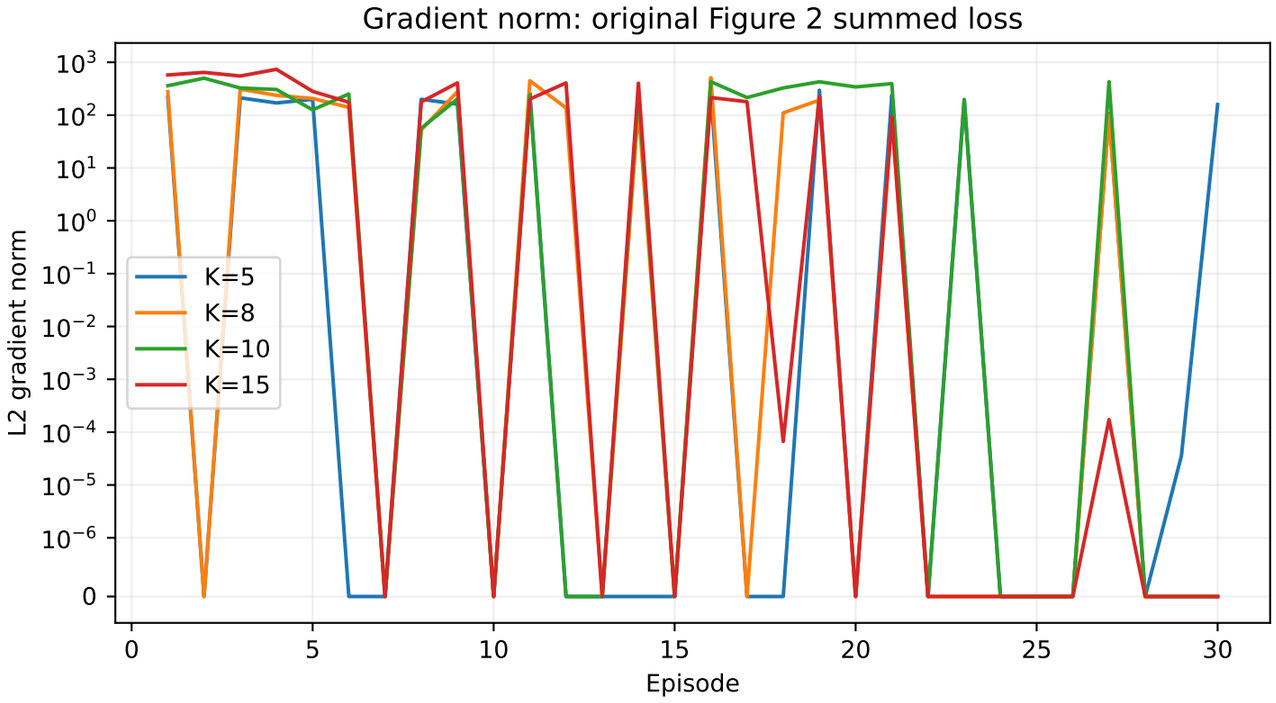

Equal-reward groups: all advantages are zero

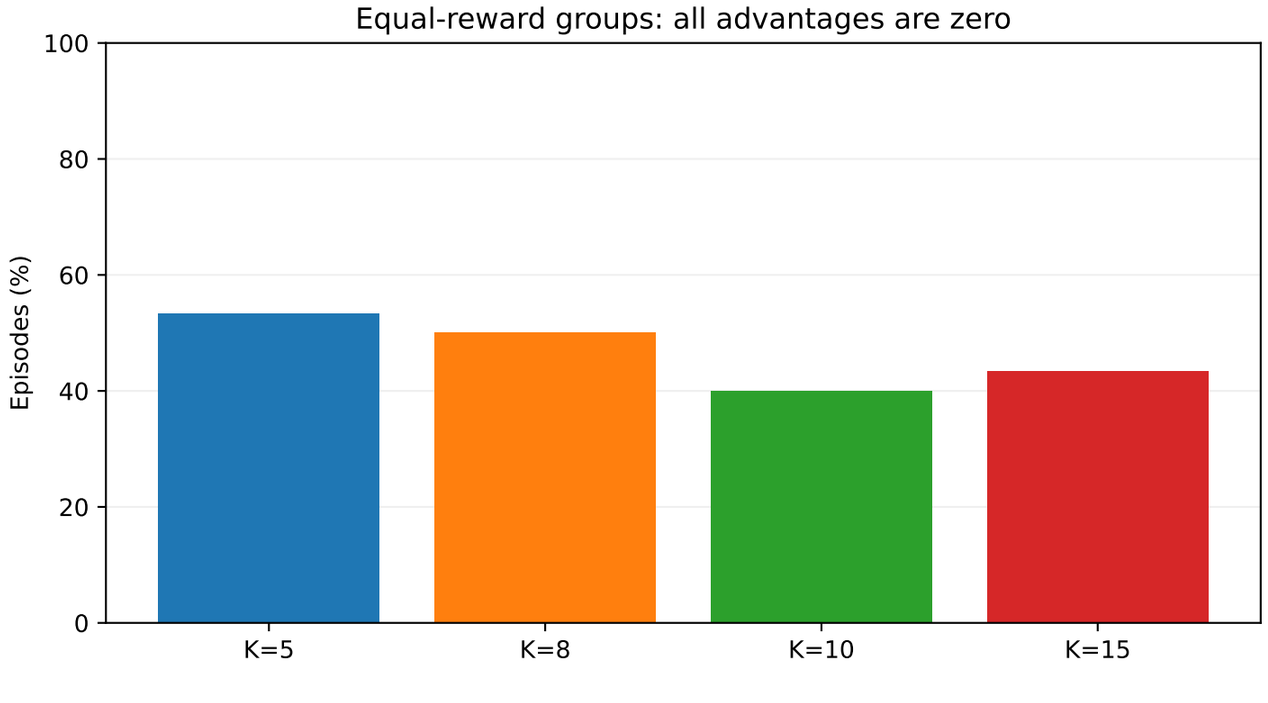

Training time, excluding validation and checkpoint I/O

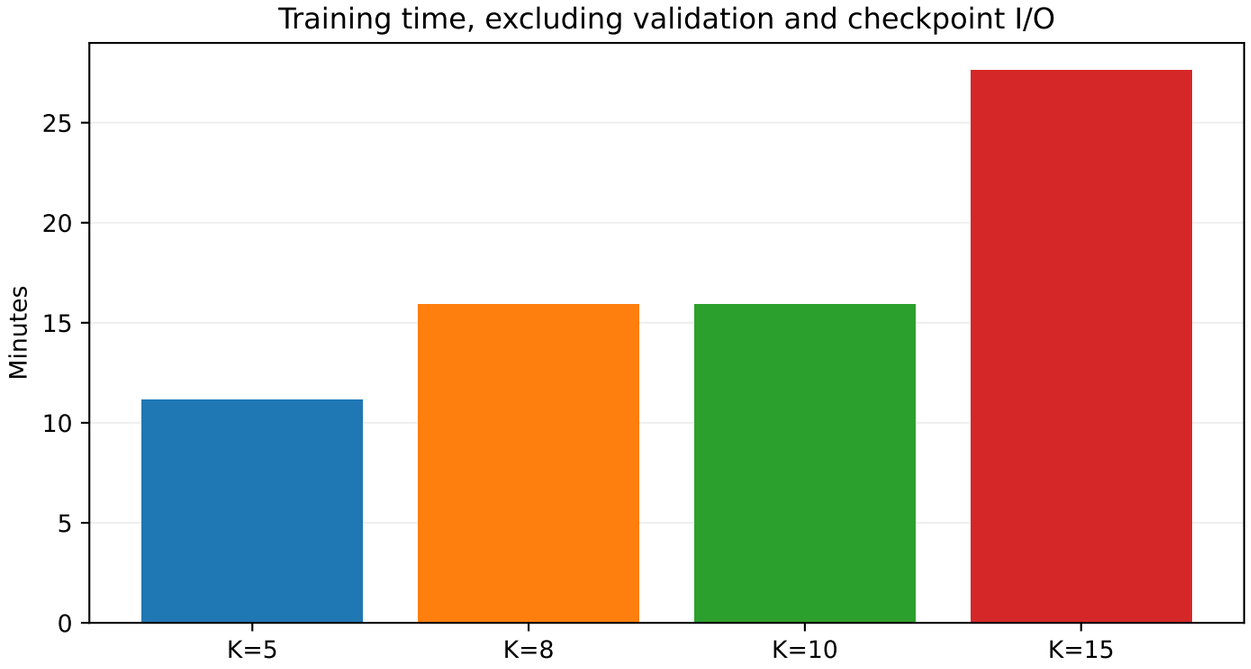

Training execution and judge cost

Source: saved run artifacts. One training seed; 30 held-out questions. Accuracy intervals describe test-set uncertainty, not variability across seeds.